

design-document

- design-document
 - Overview
 - Packaging
 - * Packaging : Overview
 - * Packaging : Diagram
 - * Packaging : Automation and Workflow Development Plan
 - * Packaging : implementation details
 - Provision
 - * Provision : Overview
 - * Provision : Diagram
 - * Provision : Automation and Workflow Development Plan
 - * Provision : implementation details
 - Deploy and Configure
 - * Deploy and Configure : Overview
 - Test
 - * Unit Tests
 - * API Tests
 - Security
 - * Security : Packaging
 - * Security : Provision
 - References

Overview

This document aims to expand and describe different stages and tasks in each step that need to be completed to deploy the Vault cluster following the SOW presented to GM Financial.




Figure 1: software delivery lifecycle

We cover the following stages (**Epics**) of the software delivery lifecycle

- **Packaging:** We include all required binaries and configuration artifacts in the form of Standard image
- **Provision:** The vault cluster requires a running environment and infrastructure. This infrastructure runs on Microsoft Azure.
- **Deploy and Configure:** Setup and configuration of various Vault auth methods and secret engines.
- **Test:** Creation of Terraform unit-tests with `terratest` framework and different API test scripts with `k6`

Packaging

Packaging : Overview

In the packaging phase, we use Hashicorp's **Packer** with **Ansible** provisioner to create a Standard Image based on **hardened RHEL**, approved by GMF's cybersecurity department. This Standard image Contains:

- Splunk Agent for log collection
- Vault Enterprise with License and associated configuration artifacts
- Bash scripts that set up DR cluster and render templated Vault configuration artifacts.
- Goss Binary and config file for confirmation and testing of the image at the end of the creation process

Packaging : Diagram

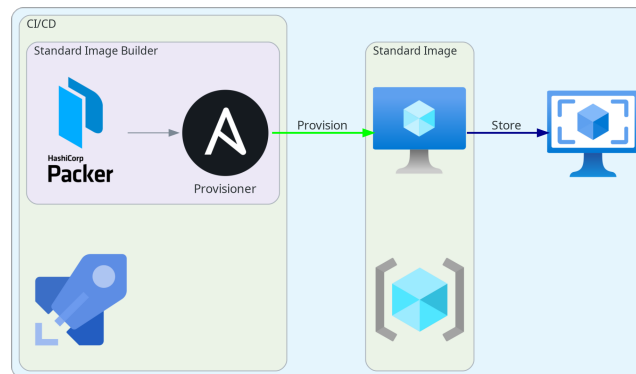


Figure 2: standard-image

Packaging : Automation and Workflow Development Plan

The following is the workflow details for developing and troubleshooting **image build** in the **Dev** environment and can be replicated in other environments:

- Have the **DevOps** team create an **Azure Compute Gallery Resource** in the **deployment** resource group.

- Have the **DevOps** team create a **Dev Image Gallery Builder** Managed Identity.
- Use a **Dev Image Gallery Builder** Managed Identity for troubleshooting builds and deployments into **Dev Azure Compute Gallery**.
- **Dev Image Gallery Builder** Managed Identity has minimal required access permission level on **Dev Azure Compute Gallery** , much like the example provided by **DevOps** team.
- The deployed Vault image needs access to Azure resources (e.g., before **vault** service starts, we would run a script that pulls tls key-pair from Azure Key Vault), so we need to have the VM Images created with a system assigned or User Assigned Managed Identity and secure their runtime access by providing access to the Managed Identity(s) to enable this access.

Packaging : implementation details

We use test-driven development (TDD) process, which leverages **goss**. The following are the requirements:

- Azure CLI must be installed and available in **PATH**.
- Ensure that Linux firewall (e.g **IPTables**) allows ingress for the following ports:
 - 8200: Vault API , Cluster bootstrapping and Vault UI
 - 8201: Raft, replication, request forwarding
 - 22: SSH connection
- Vault binary must be in user's **PATH** , under **/usr/local/bin** and ensure it is executable.
- Create **vault** system user and group and ensure it does not have a shell so no human entity can login as **vault**
- Vault Data Directory
 - Location : **/var/vault/data**
 - Ownership : **vault** user/group.
 - Stores Raft HA storage data after Vault starts
- Static Vault Configuration directory:
 - Location : **/etc/vault.d**
 - Ownership : **vault** user/group.
 - Include any 'static' (not-templated) Vault configuration artifacts.
 - Files:
 - * **/etc/vault.d/listener.hcl**: stores configuration for setting up Vault TLS listener.
 - * **/etc/vault.d/telemetry.hcl**: stores configuration that tells Vault to export metrics in **Prometheus** format.
 - * **/etc/vault.d/replication.hcl**: stores configuration parameters for tuning replication related values.
- Templated Vault Configuration Directory :
 - Location: **/var/vault.d**
 - Ownership : **vault** user/group.
 - Include templated configuration artifacts; i.e artifacts that require values that are only known after a VM is created based off of this standard image, such as Cluster node IP addresses
 - Templated artifacts have **tmpl** extension; e.g **/var/vault.d/foo.hcl.tmpl**
 - Files:
 - * **/var/vault.d/seal.hcl.tmpl**: Stores templated config for auto unseal with Azure Key Vault
 - * **/var/vault.d/raft.hcl.tmpl**: Stores templated config file for using built in Raft HA storage.
 - * **/var/vault.d/config-root.hcl.tmpl**: Stores configuration directives that do not belong to any stanzas such as **ui**
- Vault Server TLS certificates
 - Location: **/etc/vault.d/tls**

- Ownership : `vault` user/group.
- Script that renders templated Vault configuration artifacts :
 - Location: `/usr/local/bin/render-vault-config`
 - Must be executable
 - This script would use `envsubst` to render templated configuration artifacts under `/var/vault.d` and store the final result in `/etc/vault.d/`
 - After rendering templates, The following must exist in Vault config directory
 - * `/etc/vault.d/listener.hcl`
 - * `/etc/vault.d/telemetry.hcl`
 - * `/etc/vault.d/replication.hcl`
 - * `/etc/vault.d/seal.hcl`
 - * `/etc/vault.d/raft.hcl`
 - * `/etc/vault.d/config-root.hcl`
- Vault Certificate init script
 - Location: `/usr/local/bin/init-vault-tls`
 - Must be executable
 - This script would pull CA and tls key-pair from a trusted source and stored them in `/etc/vault.d/certificates`
 - After pulling the certificates, the following files must be created
 - * `/etc/vault.d/tls/ca.pem`
 - * `/etc/vault.d/tls/cert.pem`
 - * `/etc/vault.d/tls/key.pem`
- Vault DR cluster setup script
 - Location: `/usr/local/bin/vault-dr-cluster-setup`
 - Ownership : `vault` user/group.
 - This script runs after Vault daemon starts and sets up Vault DR cluster
- Use **SystemD** to manage Vault daemon:
 - Hardened unit file:
 - * Prevent Privilege escalation
 - * Disable swap
 - Uses `vault` system user/group
 - Uses **ExecStartPre** directive to run `/usr/local/bin/init-vault-tls` before running `vault`
 - Uses **ExecStartPre** directive to run `/usr/local/bin/render-vault-config` before running `vault`
 - Uses **ExecStartPost** to run `/usr/local/bin/vault-dr-cluster-setup` script after starting `vault`
 - Ensure that it is **enabled** (not **started**)
- Splunk configuration to read logs from `journald`

Provision

Provision : Overview

We use Hashicorp's **Terraform** to automate the acquisition and management of Azure infrastructure. We are assuming the following:

- An Azure **Resource Group** has already been created for our team.
- Vault Server CA and TLS key-pair has been stored in Azure Key Vault. The SystemD unit starting Vault process runs **init-vault-tls** script that we have already included in the Standard image to pull/update certificates from Azure Key Vault before it starts Vault process
- VNets for **Primary** and **DR** regions
 - **Central-US**: resources deployed in **Primary** cluster region
 - **EAST-US-2**: resources deployed in **DR** cluster region

We must also ensure required network security rules are in place:

- Allow intra-cluster communication on port 8201
- Allow incoming requests on port 8200
- Ensure all other traffic is blocked

Vault daemon will run on Linux Virtual Machine Availability Set with **five** instances. The following is Hashicorp's VM spec recommendations:

Size	Instance/VM Types	Disk Volume Specs
Small	Standard_D2s_v3,Standard_D4s_v3	1024GB(Premium_LRS)
Large	Standard_D8s_v3,Standard_D16s_v3	1024GB(Premium_LRS)

All incoming traffic from Clients to Vault goes through **Azure Load Balancer**

Provision : Diagram

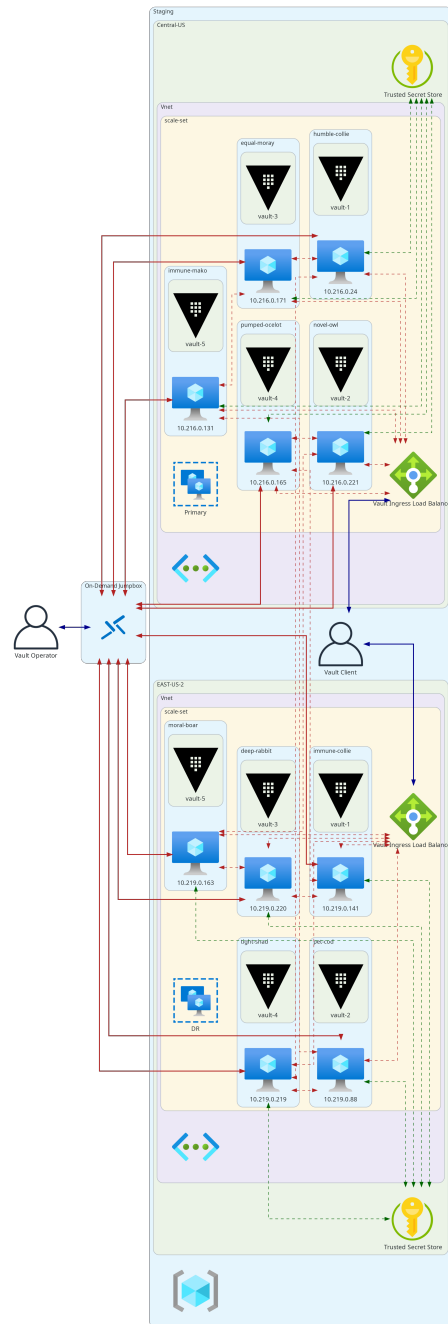


Figure 3: Infrastructure Architecture

Provision : Automation and Workflow Development Plan

The following is the workflow details for For developing and troubleshooting resource deployment in the **Dev** environment and can be replicated in other environments:

- Have the **DevOps** team create an **Azure Repo**, associated **Azure pipeline**, general access **variables group** or **variables file**, and **secrets variable group**.
- Have the **DevOps** team create a **Dev Environment Builder** Managed Identity.
- Use **Dev Environment Builder** Managed Identity for troubleshooting build and deployment from **Repo** deployment files and **Dev Azure Compute Gallery** into **Dev Resource Group**.
- Ensure **Dev Environment Builder** Managed Identity has **read-only** access to the associated **Dev Azure Compute Gallery** and permissions to create the Azure resource types needed in your deployment. Ensure the permissions are scoped only to the dedicated resource group.

Provision : implementation details

We are leveraging Hashicorp's **Terraform** to automate the acquisition and management of resources. This phase is broken down into multiple loosely connected terraform submodules. These submodules independently do not accomplish any business goals; the **root** module composes them to achieve the overall goal of provisioning the resources.

The following is a list of those submodules:

- **common-tags**: This is an auxiliary module that generates resource tags used in the other terraform submodules. It enforces a consistent naming convention for tagging resources.
- **IAM** : All Vault-related IAM resources are managed in this submodule
- **classic-load-balancer**: This submodule creates and manages a Classic load balancer for the Vault cluster.
- **vm** : This submodule creates and manages availability set nodes that Vault cluster runs on.

Deploy and Configure

Deploy and Configure : Overview

In this phase, the following are enabled and/or configured:

- Vault Namespaces
- Auth Methods:
 - App Role : Authentication method, primarily used by non-human entities.
 - JWT/OIDC: Enables authentication with JWT or OIDC provider.
 - Azure: Enables authentication against Vault using Azure Active Directory credentials
- Secret Engines:
 - KV : Generic Key-Value pair storage
 - Azure : Enables dynamic generation and lifecycle management of Azure service principals along with role and group assignments
 - PKI: Dynamic generation and management of X.509 certificates

Test

Unit Tests

Unit tests are defined based on Deployed engines. they also cover **provisioning** stage **goss** runs validation and testing at the end of the **packaging** phase

API Tests

There are many different types of API load tests. The following is a summary of them:

- **Load Testing**: type of Performance Testing used to determine a system's behavior under normal and peak conditions. The goal is to find out about:
 - assess the throughput and latency of a system under average, and peak conditions
 - Figure out SLA/SLOs
 - Ensure SLAs/SLOs are met as the system changes
- **Stress Testing**: a type of **load testing** used to determine the system's limits under extreme conditions. The goal is to find out about:
 - System breaking point and resource utilization at that point
 - Maximum acceptable request throughput before breaking
 - Recovery time of the system (if it can recover at all)
- **Spike Testing**: A type of **Stress test** that sends extreme load over a brief period. It can also be classified under **chaos experiments**. The goal is to learn about:
 - system behavior when sudden bursts happen. An example can be black Friday when retail systems are under massive load.
 - Recovery time/process when the load dies down and resource allocation patterns.
- **Soak Testing**: Type of test that focuses on system reliability. Think of this as a load test that runs continuously for a reasonably long time. Some example reliability issues include:
 - Issues due to runtime bugs
 - Memory leaks
 - Insufficient storage quotas,
 - Incorrect/poor configuration
 - infrastructure failures

Security

Security : Packaging

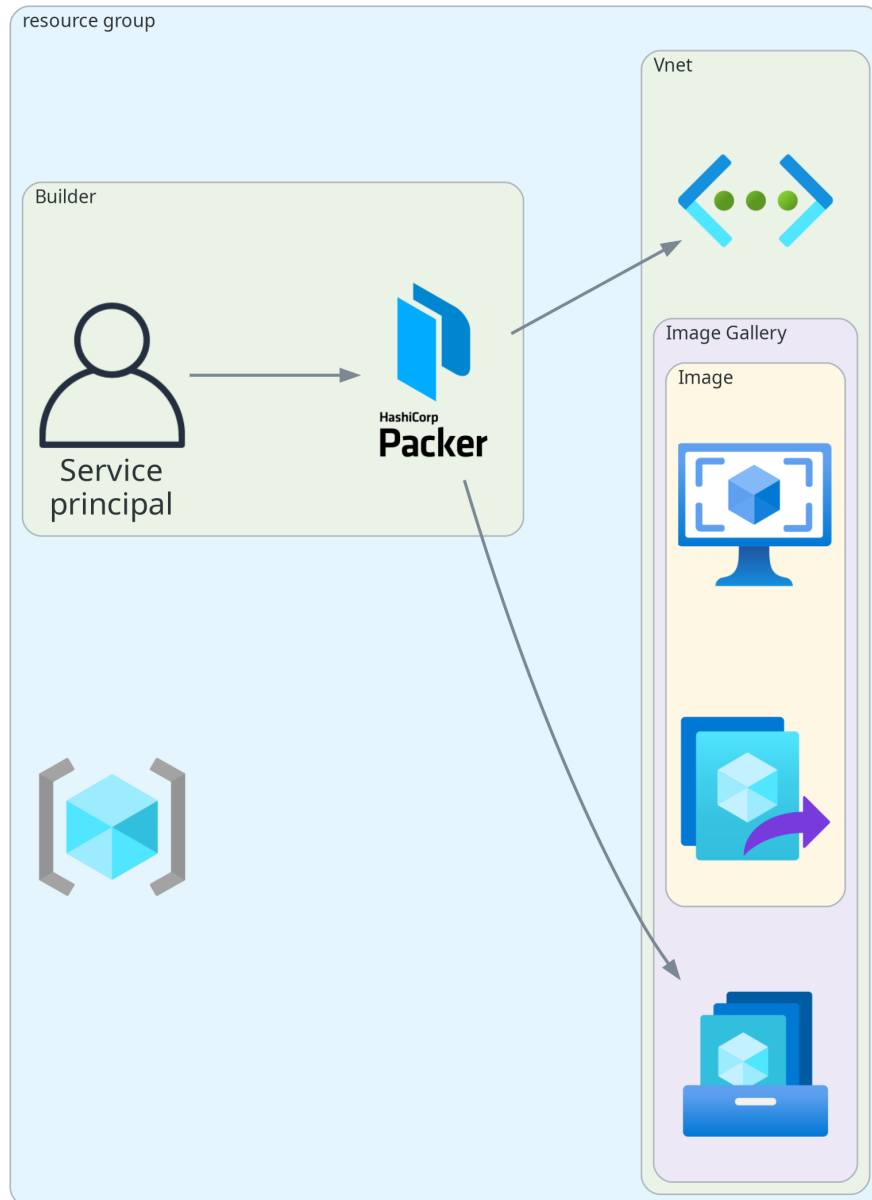


Figure 4: service principal scope

Our Vault image is based on hardened CIS Images. All operations that interact with cloud resources have to go through Azure DevOps pipelines. **Packer** uses **Azure Active Directory Service Principal**, and the following Azure **Actions** permissions in the **distribution resource group** is granted:

- Permission to customize existing images
 - Microsoft.Compute/galleries/read
 - Microsoft.Compute/galleries/images/read
 - Microsoft.Compute/galleries/images/versions/read
- Allow VM Image Builder to distribute images **within** the **distribution resource group**:
 - Microsoft.Compute/images/write
 - Microsoft.Compute/images/read
 - Microsoft.Compute/images/delete
- Allow distribution to Azure Compute Gallery:
 - Microsoft.Compute/galleries/images/versions/write
- Permission to customize images on your virtual networks
 - Microsoft.Network/virtualNetworks/read
 - Microsoft.Network/virtualNetworks/subnets/join/action

Security : Provision

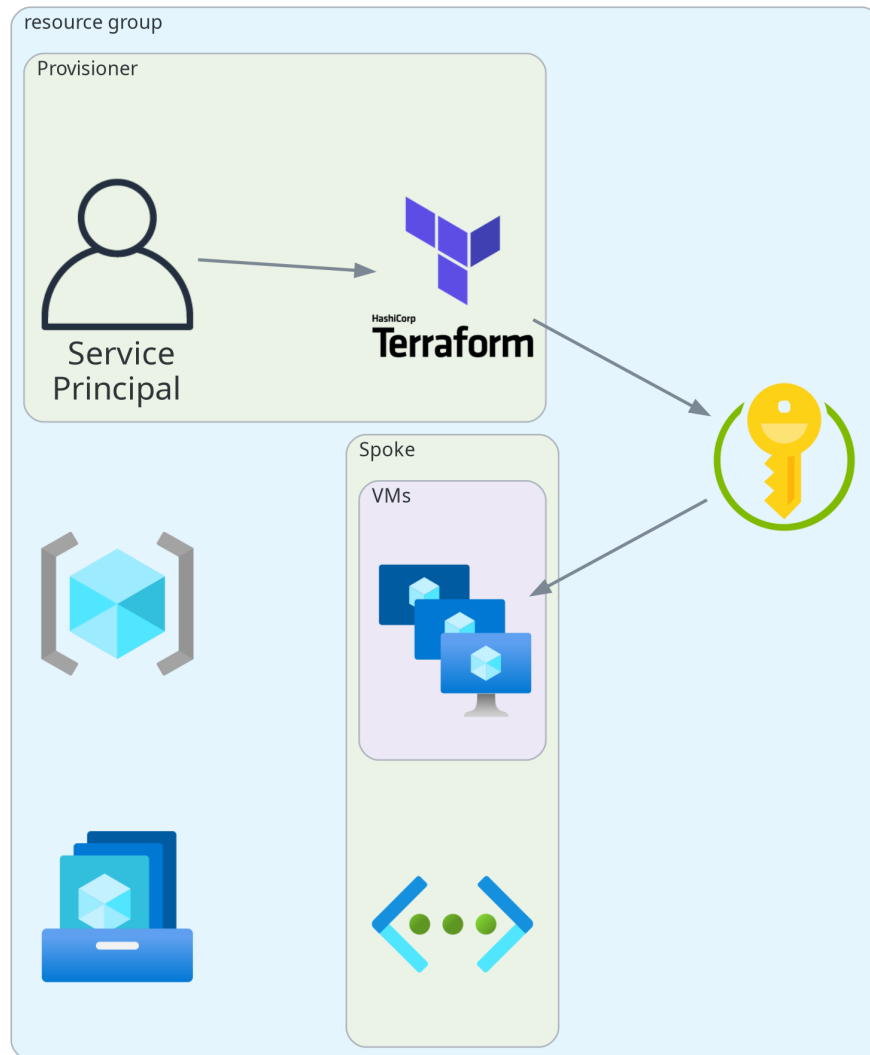


Figure 5: service principle scope

All operations that interact with cloud resources have to go through Azure DevOps pipelines.

Terraform uses Azure Active Directory Service Principal, and the following Azure Actions permissions in the distribution resource group are granted:

- resource group RBAC that allows for the creation of VMs while blocking VNet creation

```
{
  "permissions": [
    {
      "actions": [
        "*",
        "Microsoft.Compute/virtualMachines/*",
        "Microsoft.Compute/virtualMachineScaleSets/*"
      ],
      "dataActions": [],
      "notActions": [
        "Microsoft.Authorization/*/Delete",
        "Microsoft.Authorization/*/Write",
        "Microsoft.Authorization/elevateAccess/Action",
        "Microsoft.Network/dnsZones/write",
        "Microsoft.Network/dnsZones/delete",
        "Microsoft.Network/dnsZones/*/write",
        "Microsoft.Network/dnsZones/*/delete",
        "Microsoft.Network/virtualNetworks/write",
        "Microsoft.Network/virtualNetworks/delete",
        "Microsoft.Network/virtualNetworks/peer/action",
        "Microsoft.Resources/subscriptions/resourceGroups/write",
        "Microsoft.Resources/subscriptions/resourceGroups/delete"
      ],
      "notDataActions": []
    }
  ]
}
```

- The VNet needs to have an RBAC with the following permissions to allow VMs to join

```
{
  "permissions": [
    {
      "actions": [
        "Microsoft.Network/publicIPAddresses/join/action",
        "Microsoft.Network/virtualNetworks/subnets/join/action",
        "Microsoft.Network/virtualNetworks/subnets/write",
        "Microsoft.Network/virtualNetworks/*/join/action",
        "Microsoft.Network/networkSecurityGroups/write",
        "Microsoft.Network/networkSecurityGroups/securityRules/write",
        "Microsoft.Network/networkSecurityGroups/securityRules/delete"
      ],
      "dataActions": [],
    }
  ]
}
```

```

    "notActions": [],
    "notDataActions": []
  }
]
}

```

- The VMs also need to be able to read from Azure Key Vault. We can use Key Vault Secrets User built-in role with 4633458b-17de-408a-b874-0445c86b69e6 ID.
- For Vault To perform auto-unseal, it requires the following built-in roles:

Built-in role name	Role ID
Key Vault Secrets User	4633458b-17de-408a-b874-0445c86b69e6
Key Vault Crypto User	12338af0-0e69-4776-bea7-57ae8d297424

References

- [Packer Authentication for Azure](#)
- [Configure Azure VM Image Builder permissions by using the Azure CLI](#)
- [Azure built-in roles](#)
- [Create an Azure Image Builder template](#)
- [Azure Image Builder Service Image Creation Role JSON Azure security baseline for Virtual Machine Scale Sets](#)
- [Azure built-in roles for Key Vault data plane operations](#)
- [Hashicorp Learn: Auto-unseal using Azure Key Vault](#)